

Genesys Cloud Auditor — Product Roadmap

This roadmap defines the evolution of **Genesys Cloud Auditor** from a configuration audit tool into a **correlation-driven health, integrity, and escalation workbench** for Genesys Cloud organizations.

The application's purpose is to:

- detect tenant misconfigurations, inconsistencies, and hygiene issues
- identify probable platform synchronization or integrity defects
- correlate related data across multiple Genesys Cloud API domains
- produce actionable remediation guidance for internal change management teams
- generate evidence-rich escalation packets for **Genesys Care** when findings appear platform-side rather than customer-side

This roadmap builds on the current implemented checks for extension integrity, DID mismatches, queue/group hygiene, stale flows, inactive-user signals, and exported audit/operational/outbound events.

Product Vision

Genesys Cloud Auditor should answer five increasingly valuable questions:

1. **What is wrong right now?**
2. **What APIs disagree about the same object or relationship?**
3. **What changed recently that may explain the issue?**
4. **Who should act on this finding, and what should they do next?**
5. **Does this look like a tenant-side misconfiguration or a probable Genesys platform problem?**

The long-term goal is not just to export findings, but to provide a structured evidence chain that makes operational triage and support escalation dramatically faster and more reliable.

Guiding Principles

- **Correlation over isolated checks**
The highest-value findings come from comparing multiple authoritative API surfaces that should describe the same reality.
- **Actionability over raw output**
Every major finding should include severity, confidence, blast radius, recommended owner, and next step.
- **Evidence over suspicion**
Probable platform defects should only be flagged when the application can demonstrate contradictory or unstable state across trusted APIs.
- **Repeatability over one-time snapshots**
The app should preserve normalized historical snapshots so teams can distinguish chronic issues

from transient anomalies.

- **Operator clarity over mystery meat**

Findings should explain why they matter in plain language and provide object relationships, not just IDs and raw fields.

Current Implemented Capabilities

The current product already provides an effective foundation for the roadmap:

Capability	Current State
Extension consistency checks	Implemented
DID mismatch checks	Implemented
Empty / duplicate queue checks	Implemented
Empty / small group checks	Implemented
Stale flow checks	Implemented
Inactive / stale token user checks	Implemented
Missing location checks	Implemented
Audit log export	Implemented
Operational event export	Implemented
Outbound event export	Implemented
Excel workbook export	Implemented
Desktop and scheduled/headless execution model	Implemented

Roadmap Structure

Future work is grouped into five major capability layers:

1. **Configuration Integrity**
2. **Cross-Endpoint Correlation**
3. **Operational and Temporal Intelligence**
4. **Actionability and Escalation**
5. **Reporting, History, and UX**

Phase 1 — Correlation Foundation

This phase delivers the biggest leap in real-world usefulness.

1.1 Canonical correlation engine

Build a reusable engine that allows audit modules to define:

- the primary object under inspection
- related API entities to collect
- normalization rules
- expected relationships
- contradiction / anomaly rules
- severity and confidence rules
- exportable evidence fragments

Outcome

A common framework for checks that compare multiple APIs instead of one API at a time.

1.2 User–station–extension–DID–site integrity audit

Correlate:

- users
- user profiles
- assigned extensions
- DID assignments
- stations
- sites
- locations
- telephony ownership relationships

Planned checks

Check	Priority	Description
User telephony completeness	High	User appears active and telephony-enabled but lacks a coherent station/extension/DID path
DID ownership mismatch	High	A DID is associated to a user profile but assigned elsewhere in telephony
Station–user assignment conflict	High	A station references a user, but user/profile/telephony state disagrees
User–site telephony contradiction	Medium	User location/site and telephony resource site do not align
Ghost telephony assignment	High	Telephony asset references a deleted, inactive, or otherwise invalid user/resource
Multiple ownership contradiction	High	Same telephony asset appears attributable to more than one active identity

Why this matters

This is one of the clearest areas where tenant misconfiguration and platform sync issues blur together. Cross-endpoint comparison is essential.

1.3 Queue serviceability audit

Correlate:

- queues
- queue memberships
- user active state
- skills
- languages
- divisions
- wrap-up code dependencies
- routing configuration

Planned checks

Check	Priority	Description
Queue with non-serviceable membership	High	Queue has members on paper, but no realistically serviceable agents after filters are applied
Queue skill mismatch	High	Queue requires skills/languages that none of its members possess
Queue membership drift	Medium	Queue membership exists, but many members are inactive, unlicensed, or otherwise non-usable
Duplicate semantic queues	Medium	Queues are near-duplicates by name, membership, and routing intent
Queue with incomplete routing dependencies	High	Queue is missing downstream config required to actually process interactions

Why this matters

Many tenants look healthy in isolated admin screens while being operationally hollow.

1.4 Flow dependency audit

Correlate:

- Architect flows
- flow publish status
- inbound numbers / IVRs
- schedules / schedule groups

- queues referenced by flows
- downstream routing destinations

Planned checks

Check	Priority	Description
IVR–flow binding stale	High	Number points to a flow that is deleted, stale, unpublished, or errored
Dead route dependency	High	Flow references queue/schedule/destination that no longer exists
Flow dependency drift	Medium	Flow is published, but dependent objects have changed materially since last publish
Critical entry-point fragility	Medium	Important numbers depend on brittle or stale routing chains

Why this matters

This catches silent customer-facing failures before they produce obvious outage reports.

1.5 Site–edge–trunk topology integrity

Correlate:

- sites
- edges
- trunks
- stations
- telephony locations
- DID inventories where available

Planned checks

Check	Priority	Description
Trunk–edge assignment orphan	Medium	Edge references trunk/site state that no longer reconciles
Site–edge mismatch	High	Site relationship differs across authoritative edge/site resources
Station topology contradiction	Medium	Station/site/location relationships appear internally inconsistent
DID inventory orphan	Medium	DID ranges exist in inventory but do not reconcile to active service paths

Phase 2 — Operational and Time-Based Intelligence

This phase adds “what changed?” and “is this recurring?”—the good stuff.

2.1 Change-to-symptom correlation

Correlate current findings with:

- audit logs
- operational events
- outbound events
- relevant configuration changes
- flow publish timestamps

Planned checks

Check	Priority	Description
Regression chain builder	High	Build a timeline showing config changes followed by symptoms/findings
Change adjacency marker	High	Surface recent changes that touched related objects within the finding window
Suspect release window	Medium	Group findings by likely onset window after correlated changes

Example output

“Queue membership changed at 14:03; routing anomalies began at 14:07; serviceability fell to zero by 14:11.”

2.2 Flapping and instability detection

Planned checks

Check	Priority	Description
Assignment flapping	High	Ownership or assignment repeatedly changes between states
Publish churn	Medium	Flows repeatedly republished or altered without stabilizing behavior
Resource oscillation	Medium	Site/edge/station/trunk relationships repeatedly move between valid and invalid interpretations

Why this matters

Flapping is often a signal of automation conflict, admin collision, sync lag, or deeper platform instability.

2.3 Hot spot ranking

Planned checks

Check	Priority	Description
Chronic object ranking	Medium	Rank queues, users, sites, and telephony resources by repeated anomaly association
Domain instability index	Medium	Score routing, telephony, identity, or outbound domains for recurring problems
Blast-radius concentration	Medium	Identify small sets of objects involved in a disproportionate share of findings

Phase 3 — Actionability and Escalation Intelligence

This phase makes the tool genuinely operational.

3.1 Recommended action model

Every major finding should emit:

- finding type
- severity
- confidence
- blast radius
- impacted objects
- suspected owner
- recommended next action
- probable cause category

Action categories

Category	Description
Local Configuration Fix	Likely resolvable directly by tenant admin or engineering team
Change Review Required	Recent changes suggest managed review or rollback may be needed
Monitor / Re-run	Evidence is weak or transient; continue observation
Escalate to Genesys Care	Evidence suggests probable platform-side issue or unresolved contradictory state

3.2 Probable platform issue qualification

A finding should be classified as **Probable Platform Issue** only when criteria are met.

Qualification rules

A finding is a candidate for support escalation when:

1. two or more authoritative APIs disagree
2. the contradiction persists across repeated collection or repeated runs
3. no recent admin change plausibly explains the issue
4. the inconsistency reflects an invalid or impossible state under the documented model
5. business impact is meaningful or plausibly emerging

Planned checks

Check	Priority	Description
Persistent contradiction detector	High	Contradictory API state persists across collections/runs
No-local-cause qualifier	High	Finding remains after excluding recent tenant changes as an explanation
Support-readiness scorer	High	Finding contains sufficient object IDs, timestamps, evidence, and impact context for case creation

3.3 Genesys Care evidence packet export

Generate a structured escalation packet for support-worthy findings.

Packet contents

- org / tenant identifier
- region
- generated timestamp
- affected object IDs and names
- API surfaces involved
- contradictory evidence summary
- related audit log / operational event timeline
- reproduction notes if available
- severity, confidence, and blast radius
- recommended case summary text
- supporting workbook sheet references

Planned outputs

Output	Priority	Description
Support case summary	High	Human-readable case narrative for Care ticket entry

Output	Priority	Description
Evidence JSON	High	Machine-readable evidence package for archival or automation
Timeline appendix	Medium	Event and change history supporting the escalation

Phase 4 — Historical Baselines and Drift Intelligence

This phase gives the application memory.

4.1 Snapshot persistence

Persist normalized snapshots for each audit run:

- normalized users
- normalized telephony ownership
- normalized queues and memberships
- normalized flow dependencies
- normalized topology relationships
- finding summaries

Planned features

Feature	Priority	Description
Snapshot save/load	High	Save normalized state for comparison across runs
Historical diff engine	High	Compare current state to previous baseline(s)
Chronic anomaly ledger	Medium	Track findings that recur over time
Finding lifecycle classification	Medium	Mark findings as New, Resolved, Recurrent, or Worsening

4.2 Drift detection

Planned checks

Check	Priority	Description
Membership drift	Medium	Queue/group/user relationships changed materially since last baseline
Telephony ownership drift	High	Extensions, DIDs, stations, or site bindings changed unexpectedly
Routing drift	High	Flow/queue/schedule bindings changed between runs
Security drift	Medium	Roles, OAuth clients, or privileged assignments changed materially

Phase 5 — Reporting, UX, and Explainability

This phase turns the results into something people will actually use.

5.1 Executive summary and triage dashboard

Planned outputs

Output	Priority	Description
Executive summary sheet	High	Findings by severity, domain, probable cause, and owner
"Open case recommended" counter	High	Count and summarize support-worthy findings
Top impacted objects	Medium	Highlight objects with highest recurrence or blast radius
Domain health scoring	Medium	Score telephony, routing, identity, outbound, and security posture

5.2 Relationship and dependency views

Planned outputs

Output	Priority	Description
Object dependency tree	High	Render finding-specific object relationship chains
Finding evidence chain	High	Show which APIs and comparisons produced the conclusion
Why-this-matters explanation	High	Plain-language explanation of risk or likely impact
Recent change context	Medium	Inline summary of nearby relevant changes/events

Example relationship chain

User -> Extension -> DID -> Station -> Site -> Queue -> Flow

or

DID -> IVR -> Flow -> Queue -> Agents

5.3 Workbook and export improvements

Planned features

Feature	Priority	Description
---------	----------	-------------

Feature	Priority	Description
Summary-first workbook layout	High	Open directly to triage-oriented summary sheets
Cross-sheet linking	Medium	Link findings to supporting detail sheets
JSON evidence export	High	Parallel machine-readable export for automation
HTML summary export	Medium	Human-friendly evidence summary suitable for sharing or ticket attachments
Rule metadata export	Medium	Include which audit rule/version produced each finding

Cross-Domain Audit Catalog

The following catalog groups future checks by domain rather than implementation phase.

Identity / Access

Planned Check	Priority	Description
Inactive users with privileged roles	High	High-impact access remains assigned to stale or inactive accounts
Roles assigned to invalid users	Medium	Roles reference deleted or invalid identities
Over-privileged OAuth clients	High	OAuth clients hold permissions broader than intended
Division scope contradictions	Medium	User/role/division relationships produce misleading or inconsistent access state
Privileged asset ownership drift	Medium	Sensitive admin or telephony ownership changed unexpectedly

Routing / Queueing

Planned Check	Priority	Description
Queue with no routing viability	High	Queue cannot effectively service work despite appearing configured
Queue dependency gap	High	Queue relies on missing wrap-up/skill/language/routing configuration
Queue duplication cluster	Medium	Multiple queues likely represent abandoned clones or overlapping intent

Planned Check	Priority	Description
Routing-policy contradiction	Medium	Queue behavior implied by one resource is contradicted elsewhere

Telephony / Numbering

Planned Check	Priority	Description
DID ownership mismatch	High	DID profile ownership disagrees with telephony assignment
Extension range violation	Medium	Extensions fall outside expected site/org numbering patterns
User-station-DID contradiction	High	Cross-object telephony identity does not reconcile
Emergency service exposure	High	Telephony-enabled entity lacks coherent emergency-related assignment context
Inventory orphan detection	Medium	Number or asset inventory exists without valid downstream binding

Architect / Flowing / Entry Points

Planned Check	Priority	Description
Stale dependency after publish	Medium	Published flow depends on changed or missing downstream resources
Number-to-flow fragility	High	Critical inbound numbers depend on stale or brittle routing
Schedule reference contradiction	Medium	Schedule or schedule group state no longer supports intended routing
Orphaned destination references	High	Flow references deleted queues, flows, or destinations

Outbound

Planned Check	Priority	Description
Campaign with invalid dependency	High	Campaign depends on inactive or deleted object such as list/queue
Outbound routing contradiction	Medium	Outbound configuration implies invalid downstream service path

Planned Check	Priority	Description
Outbound event hot spot	Medium	Specific campaigns or lists correlate disproportionately to anomalies
Change-linked outbound degradation	Medium	Recent changes align with outbound issue onset

Data Quality / Hygiene

Planned Check	Priority	Description
Duplicate semantic names	Medium	Similar names likely represent confusion or abandoned copies
Missing user email / metadata	Low	Weak identity hygiene that can affect reporting and administration
Naming standard drift	Low	Objects diverge from expected naming patterns by domain/site/division
Legacy orphan cluster	Medium	Objects appear to be remnants of prior migrations or decommissioned designs

Rule Authoring and Extensibility

To make the platform durable, future audit checks should be definable through a common rule contract.

Proposed rule contract attributes

- rule ID
- rule name
- domain
- required endpoints / collectors
- normalization dependencies
- join keys / correlation keys
- anomaly criteria
- severity default
- confidence model
- evidence template
- recommended owner
- suggested actions
- support-escalation eligible (true/false)

This enables:

- internal extensibility

- customer-specific rule packs
- controlled rule versioning
- regression-safe testing

Testing and Validation Roadmap

Planned quality gates

Gate	Purpose
Collector contract tests	Ensure endpoint ingestion remains stable
Correlation tests	Validate cross-endpoint joins and contradiction logic
Severity/confidence tests	Prevent noisy or inflated findings
Snapshot diff tests	Ensure historical change tracking is deterministic
Export validation tests	Verify workbook/JSON/HTML output integrity
Care packet fixture tests	Ensure escalation bundles contain the expected evidence set

Suggested Delivery Sequence

Near-term priority

1. Correlation engine
2. User–station–extension–DID–site integrity audit
3. Queue serviceability audit
4. Flow dependency audit
5. Action model (owner + next step)
6. Support case evidence packet export

Mid-term priority

1. Snapshot persistence
2. Drift engine
3. Change-to-symptom correlation
4. Flapping detection
5. Topology integrity checks

Long-term priority

1. Domain health scoring
2. Rule-pack extensibility
3. HTML evidence reports
4. Advanced trend analytics

Definition of Success

Genesys Cloud Auditor is successful when it can reliably produce findings that are:

- **clear enough for admins to fix**
- **credible enough for engineers to trust**
- **structured enough for automation to consume**
- **evidence-rich enough for Genesys Care escalation**
- **historically grounded enough to distinguish noise from chronic problems**

At that point, the application becomes more than an auditor:
it becomes a tenant integrity and platform anomaly investigation tool.